

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import make_scorer, mean_absolute_error, mean_squared_error
```

```
In [2]: warnings.filterwarnings('ignore')
```

```
In [3]: # Load the datasets
df_red = pd.read_csv(r"C:\Users\USER\Downloads\wine\winequality-red.csv", sep=";")
df_white = pd.read_csv(r"C:\Users\USER\Downloads\wine\winequality-white.csv", sep=";")
```

```
In [4]: # Combine the datasets
df = pd.concat([df_red, df_white], ignore_index=True)
```

```
In [5]: # Check for null values
print(df.isnull().sum())
```

```
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density           0
pH                0
sulphates         0
alcohol           0
quality           0
dtype: int64
```

```
In [6]: # Check the basic info of the dataframe
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6497 entries, 0 to 6496
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          6497 non-null   float64
1   volatile acidity       6497 non-null   float64
2   citric acid            6497 non-null   float64
3   residual sugar         6497 non-null   float64
4   chlorides              6497 non-null   float64
5   free sulfur dioxide    6497 non-null   float64
6   total sulfur dioxide   6497 non-null   float64
7   density                6497 non-null   float64
8   pH                    6497 non-null   float64
9   sulphates              6497 non-null   float64
10  alcohol                6497 non-null   float64
11  quality                6497 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 609.2 KB
None
```

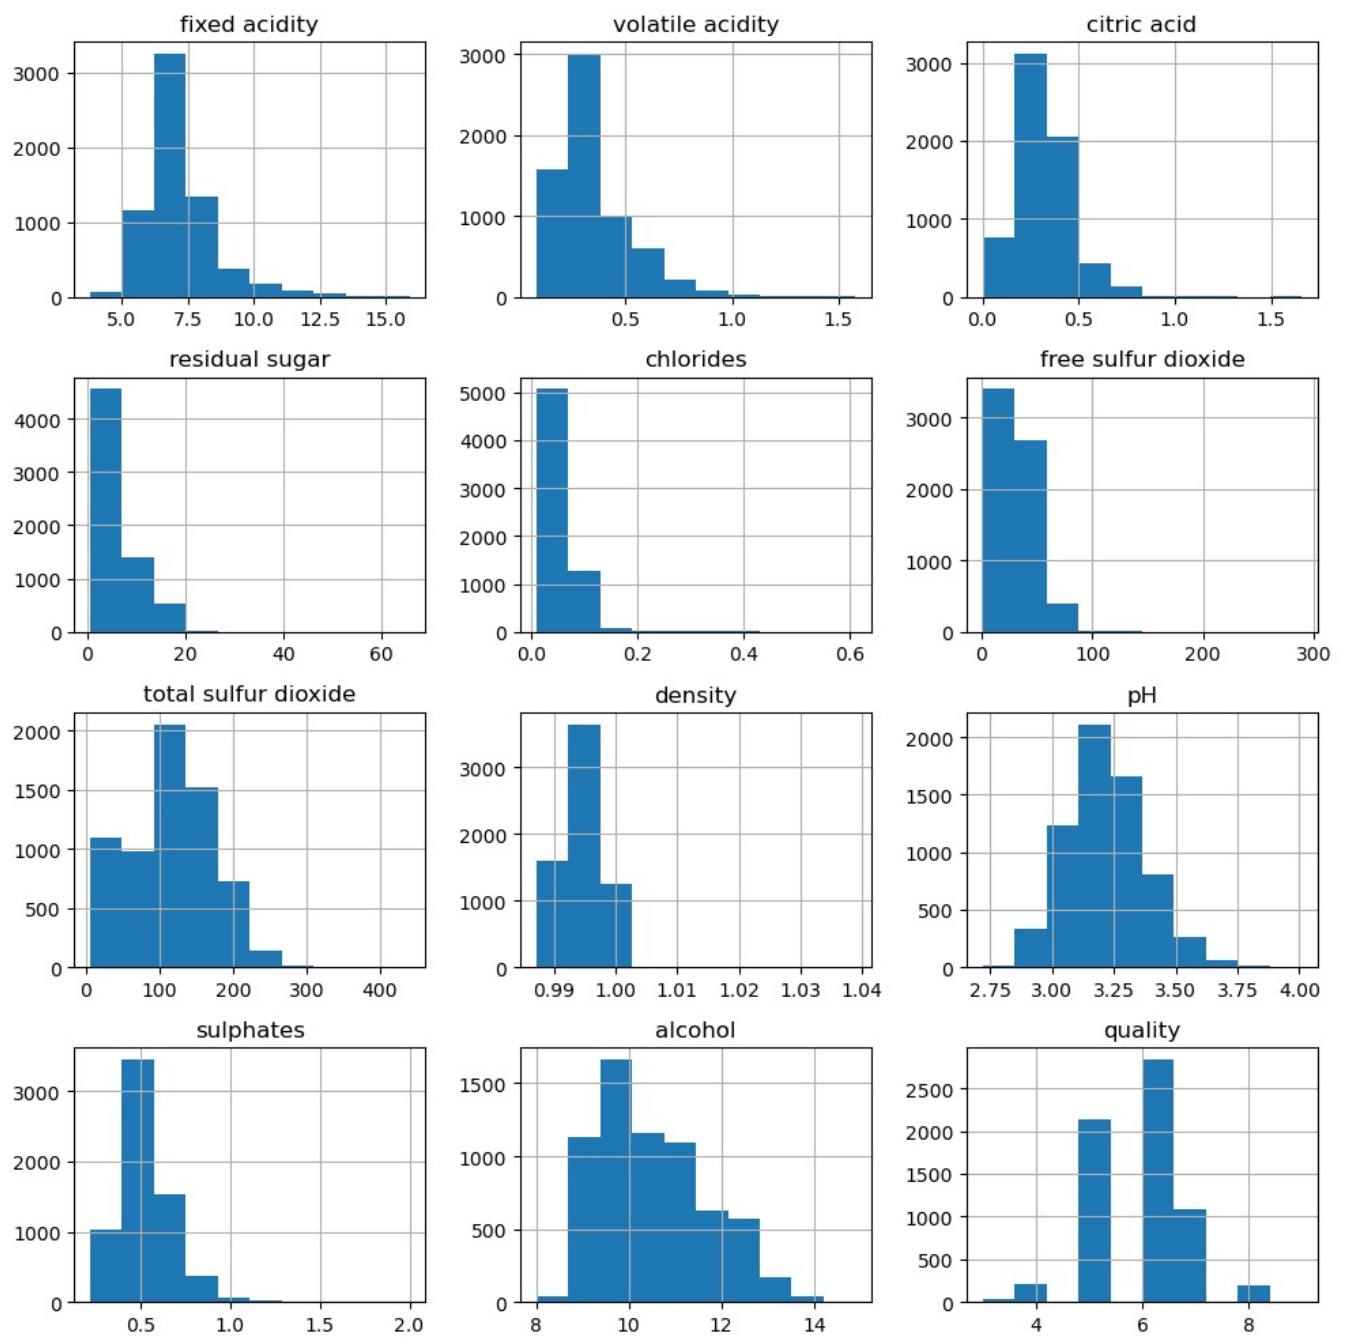
```
In [7]: # Describe the data
print(df.describe())
```

	fixed acidity	volatile acidity	citric acid	residual sugar \
count	6497.000000	6497.000000	6497.000000	6497.000000
mean	7.215307	0.339666	0.318633	5.443235
std	1.296434	0.164636	0.145318	4.757804
min	3.800000	0.080000	0.000000	0.600000
25%	6.400000	0.230000	0.250000	1.800000
50%	7.000000	0.290000	0.310000	3.000000
75%	7.700000	0.400000	0.390000	8.100000
max	15.900000	1.580000	1.660000	65.800000

	chlorides	free sulfur dioxide	total sulfur dioxide	density \
count	6497.000000	6497.000000	6497.000000	6497.000000
mean	0.056034	30.525319	115.744574	0.994697
std	0.035034	17.749400	56.521855	0.002999
min	0.009000	1.000000	6.000000	0.987110
25%	0.038000	17.000000	77.000000	0.992340
50%	0.047000	29.000000	118.000000	0.994890
75%	0.065000	41.000000	156.000000	0.996990
max	0.611000	289.000000	440.000000	1.038980


	pH	sulphates	alcohol	quality
count	6497.000000	6497.000000	6497.000000	6497.000000
mean	3.218501	0.531268	10.491801	5.818378
std	0.160787	0.148806	1.192712	0.873255
min	2.720000	0.220000	8.000000	3.000000
25%	3.110000	0.430000	9.500000	5.000000
50%	3.210000	0.510000	10.300000	6.000000
75%	3.320000	0.600000	11.300000	6.000000
max	4.010000	2.000000	14.900000	9.000000

```
In [8]: # Plot histograms for all features
df.hist(figsize=(10, 10))
plt.tight_layout()
plt.show()
```




```
In [9]: # Boxplots for each feature
for column in df.columns:
    plt.figure(figsize=(8, 6))
    plt.boxplot(df[column])
    plt.title(f'Box Plot of {column}')
    plt.show()
```


Box Plot of fixed acidity




Box Plot of volatile acidity




Box Plot of citric acid




Box Plot of residual sugar




Box Plot of chlorides




Box Plot of free sulfur dioxide




Box Plot of total sulfur dioxide




Box Plot of density





Box Plot of pH




Box Plot of sulphates







```
In [10]: # Correlation heatmap
plt.figure(figsize=(20, 10))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.show()
```



```
In [11]: # Prepare data for machine learning
X = df.drop('quality', axis=1)
y = df['quality']
```

```
In [12]: # Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=89)
```

```
In [13]: # Print the shape of the training and testing data
print(X_train.shape, X_test.shape)

(5197, 11) (1300, 11)
```

```
In [14]: # Create the pipeline
pipeline = make_pipeline(StandardScaler(), RandomForestRegressor(random_state=42))
```

```
In [15]: # Using R-squared for cross-validation
r2_scores = cross_val_score(pipeline, X, y, scoring='r2', cv=5)
mean_r2 = r2_scores.mean()
print(f'Mean R-squared: {mean_r2:.2f}')
```

Mean R-squared: 0.27

Interpretation of R-squared:

0 to 1 Range:  $R^2$  ranges from 0 to 1. An  $R^2$  value closer to 1 indicates a better fit. Low  $R^2$  (0.27): This suggests that the model explains only a small portion of the variance, implying a potentially weak model for predicting wine quality based on the given features.

```
In [16]: from sklearn.linear_model import LogisticRegression
```

```
In [17]: # Create a new pipeline with LogisticRegression
lr_pipeline = make_pipeline(StandardScaler(), LogisticRegression(random_state=42, max_iter=10000))
```

```
In [18]: lr_scores = cross_val_score(lr_pipeline, X, y, scoring='accuracy', cv=5, n_jobs=-1)
mean_lr_accuracy = lr_scores.mean()
print(f'Mean Accuracy: {mean_lr_accuracy:.2f}')
```

Mean Accuracy: 0.50

```
In [21]: from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
```


```
In [22]: rf_pipeline = make_pipeline(StandardScaler(), RandomForestClassifier(random_state=42))
```

```
In [23]: param_grid = {
    'randomforestclassifier__n_estimators': [50, 100, 200],
    'randomforestclassifier__max_depth': [None, 10, 20, 30],
    'randomforestclassifier__min_samples_split': [2, 5, 10],
    'randomforestclassifier__min_samples_leaf': [1, 2, 4]
}
```

```
In [24]: grid_search = GridSearchCV(rf_pipeline, param_grid, cv=5, scoring='accuracy', n_jobs=-1)
```

```
grid_search.fit(X_train, y_train)
```

Out[24]:



```
In [26]: # Evaluate best model from Grid Search
best_rf_model = grid_search.best_estimator_
```

```
In [27]: # Cross-validation score
rf_scores = cross_val_score(best_rf_model, X, y, scoring='accuracy', cv=5, n_jobs=-1)
mean_rf_accuracy = rf_scores.mean()
print(f'Mean Accuracy with Random Forest: {mean_rf_accuracy:.2f}')
```

Mean Accuracy with Random Forest: 0.48

Okay, here we try 3 models but, its not working properly any of the model, we may need to do more some feature engineering, we will try second version of this code again

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js